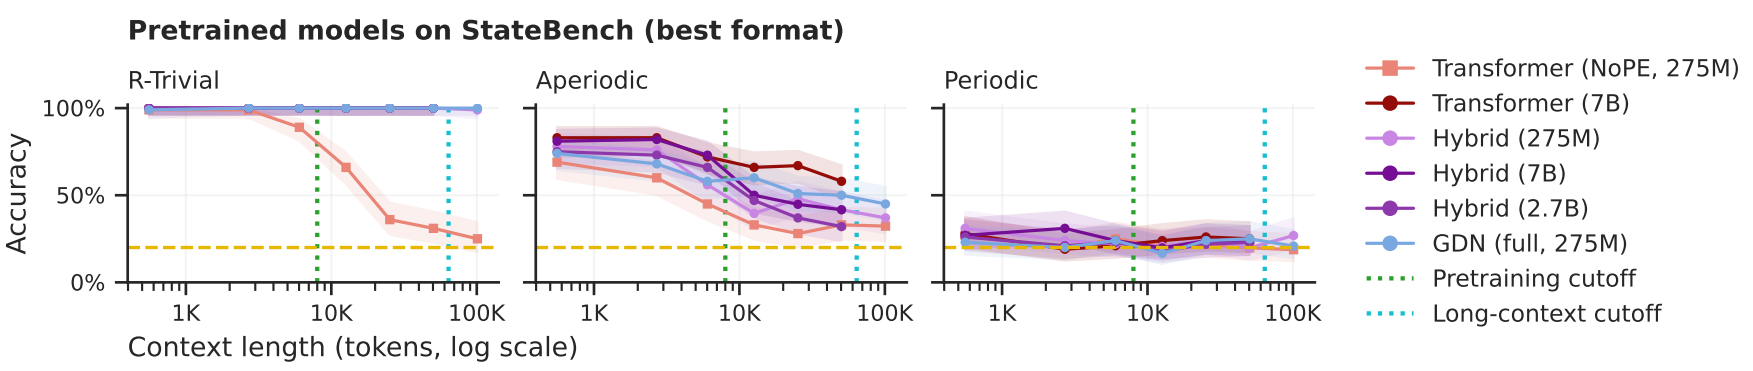
Pretrained models on StateBench (best format)
R-Trivial
Aperiodic
Periodic
Accuracy
100%
50%
0%
1K
10K
100K
Context length (tokens, log scale)
Transformer (NoPE, 275M)
Transformer (7B)
Hybrid (275M)
Hybrid (7B)
Hybrid (2.7B)
GDN (full, 275M)
Pretraining cutoff
Long-context cutoff