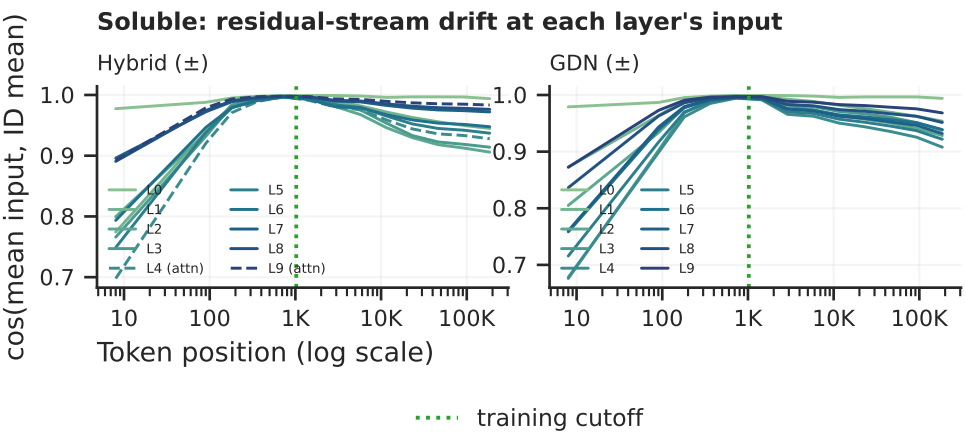
Soluble: residual-stream drift at each layer's input
Hybrid (±)
GDN (±)
cos(mean input, ID mean)
1.0
0.9
0.8
0.7
L0
L1
L2
L3
L4 (attn)
L5
L6
L7
L8
L9 (attn)
L4
L9
10
100
1K
10K
100K
Token position (log scale)
training cutoff